

SciOdyssey.

A research layer over your agent harness.

Our agent can carry a research task to completion.

PRESERVE THE RESEARCH WORLD.
RESET THE RESEARCHER.

SAME
EVIDENCE

FRESH
JUDGMENT
A WIDER
SEARCH

Research
Runs
Longer.

01 / PROBLEM

What sustained research needs.


Open-ended research changes the question while it runs. A single reasoning trail can anchor the search, repeated context is costly, and a valid result needs evidence that survives the next attempt.

ADAPT RESET REVIEW DELIVER



02 / SYSTEM

One persistent world. Fresh scientific judgment.



Audit → scope → preserve → fresh Scientist
Selective persistence keeps one implementation parent and a traceable evidence record.

HARNESS + RUNTIME workspace isolation · runner + evaluator · logging · artifact capture · State operations

Parallel breadth isolated branches → post-hoc review → explicit adoption

PERSISTENT RESEARCH WORLD

What survives the next trajectory

VERIFIED / EVIDENCE	CURATED / FALLIBLE	IMPLEMENTATION
Knowledge + experiment ledger Metrics, failures, original reports	Research brief + priors Current understanding, open to revision	Retained code + State Recoverable version with provenance

The next Scientist inherits experience, not the previous Scientist's reasoning.

03 / EVALUATION

Five claims. Evidence in the card.

KuaiRand-Pure · fixed evaluator · public validation · recorded run

01 MEASURED RESULT

The model got better.

0.6016000 → 0.6059363

official FM SciOdyssey · +0.0043363

GAUC nDCG@5 46 fields · 8 seeds · CPU

0.6728421 0.5390304

02 SUSTAINED SEARCH

The search kept going.

0.6016310 0.6040901 0.6044289 0.6059363

C1 baseline C2 expand C3 variance C4 retained

4 cycles · 13 named experiments · 7 Full evaluations

03 RECOVERY UNDER FAILURE

Our agent is robust.

FAIL REPAIR RETAIN

NumPy scalar breaks writer → Scientist fixes serialization → rerun keeps evidence

Training and evaluation had completed; the failure happened at evidence writing. META reviewed what became durable State.

04 DIRECT CODEX CONTROL

Our agent beat Codex on this task.

0.6044533 → 0.6059363

direct Codex Goal SciOdyssey
provisional public validation

Δ +0.0014830 Primary. Different systems and protocols; this is a recorded task result, not a causal or general benchmark claim.

05 RESOURCE ACCOUNTING

We are not tokenmaxxing.

~US\$10 subscription-equivalent project usage GPT Plus + Gemini Pro · 7 days
1 × MacBook M2

48.240M 4.020M 0

total tokens including cache-read non-cache input + output GPU-hours NumPy / CPU

Public validation only. Hidden-test scoring is organizer-controlled. Resource figures use their stated project-level and retained-run scopes.

04 / INSIGHTS

Observed behavior, not just a score.

RECOVERY Evidence failure became a repairable event instead of a dead end.	AUDIT META scopes claims against artifacts; it is not a formal verifier.	OPEN QUESTION One task does not establish generality. Reset and State effects remain coupled.
--	--	---

05 / EXPERIENCE

Start your journey.

task.md

What to solve

Objectives · constraints · evaluation

PERSONAL.md

How to work

Preferences · research style · priorities

Define → Step → Run → Parallel → Inspect

Autonomous by default.
Supervisable by design.

○ ○ ○ Coding agent

EXAMPLE INPUT

Use the research-agent skill.
Follow task.md and PERSONAL.md.
Let me review the results and evidence.