

SciOdyssey.

A research layer over your agent harness.

Our agent can carry a research task to **completion.**

— PRESERVE THE RESEARCH WORLD · RESET THE RESEARCHER

ROBUST

SUPERVISABLE

BROAD SEARCH

SAME EVIDENCE
FRESH JUDGMENT
A WIDER SEARCH

Research
Runs
Longer.

01 / PROBLEM —

Why current agents fall short.

Sustained autonomous exploration encounters three critical bottlenecks:

01 Closed-world fragility

Rigid predefined pipelines shatter as soon as unexpected runtime errors emerge. In Cycle 1, a **float32 JSON serialization bug** halted logging; **without autonomous self-healing, exploration dies permanently.**

02 Single-trajectory inertia

Accumulating context history anchors the agent to past reasoning, **trapping search in local basins** (e.g., standard single-agent baselines **stall at 0.6044** while endlessly fine-tuning hyperparameters).

03 Context inflation cost

Continually rereading deep conversational context **costs 8x more tokens** than the actual code edits, exhausting context windows and causing **severe reasoning dilution.**

TRADITIONAL AGENT HARNESS

Single Monolithic Context · Error Cascades · Token Waste · Local Optima

SCIODYSSEY RESEARCH LAYER

Ephemeral Fresh Scientist · Persistent Disk Memory · Verified Handoff

PARADIGM SHIFT

Decouple the **persistent research world** (disk memory) from ephemeral **researcher reasoning** (fresh context).


02 / SYSTEM —

One persistent world. Fresh scientific judgment.

Give the research world and the researcher different lifetimes.


THE DUAL-LIFETIME ENGINE

Context reset between trajectories · Audited persistence on disk



PERSISTENT RESEARCH WORLD

What survives across trajectories



The next Scientist inherits **verified experience, not prior reasoning inertia.**

Decoupled Memory · Audited Persistence · Autonomous Self-Healing · Zero Context Pollution

03 / EVALUATION —

Five empirical claims. Measured evidence.

Standardized evaluation on KuaiRand-Pure benchmark · Fixed official evaluator · Complete telemetry accounting

01 MEASURED RESULT

The model got better.

0.6016000

Official FM baseline



0.6059363

SciOdyssey submission

+0.0043363 Primary improvement

GAUC 0.6728421

(+0.0054)

nDCG@5 0.5390304

(+0.0033)

LogLoss 0.38102

(-0.0032)

Feature Engine: 46 categorical cross-interactions + log1p transforms

Architecture: 8-seed Factorization Machine ensemble (NumPy / CPU)

Monotonic gain from baseline (0.6016) → feature engineering (0.6041) → seed ensembling (0.6059); **100% CPU inference** without GPU clusters.

02 SUSTAINED SEARCH

The search stayed productive.

0.60163

C1: Valid Base

0.60409

C2: Features

0.60443

C3: Var. Red.

0.60594

C4: 8-Seed FM

SciOdyssey	0.6059363	Climbed across 4 cycles (7 Full evals)
Codex Goal	0.6044533	Peaked at C8; stalled & dropped in 22 runs

4 autonomous cycles

7 Full evaluations

+0.00148 over Codex

7 Full public-validation evaluations retained across 4 cycles. Trajectory climbed via **structural feature expansion** rather than single lucky edits.

03 SELF-HEALING & GATEKEEPING

The agent recovered from failure.

TIER 1 · SCIENTIST SELF-HEALING Cycle 1 runtime
Traceback auto-repaired: numpy.float32 broke JSON serialization. Scientist diagnosed traceback, patched scalar serialization helper, and reran cleanly.

TIER 2 · META GATEKEEPING Proxy audit receipt
Spurious proxy rejected: Scientist claimed false gain; Meta audited system/data.py, caught UNK user sampling distortion, and blocked State promotion.

Self-Healing: 0 human bugfixes


META Gate: Spurious proxy killed

Audit: 170,588 rows aligned

Robust research requires both **execution-level crash recovery** and **epistemic defense** against false breakthroughs. **Zero human debugging interventions across 7 days.**

05 DIRECT AGENT BENCHMARK

Score versus measured LLM-token investment.



All results **verified on KuaiRand-Pure official validation split** with **complete reproducible artifacts.**

04 / INSIGHTS —

What the run taught us.

01 · MODEL DIVERSITY

Depth + Breadth Synergy.

GPT 5.6 Luna (depth: refine ideas) + **Gemini 3.7 Flash** (breadth: open new directions) in one shared world pushed score from **0.6052** (Gemini-only) to **0.6059363** (GPT+Gemini).

Gemini-only 0.6052 → GPT+Gemini 0.6059363

02 · PARALLEL SYNTHESIS

Cross-Branch Idea Transfer.

Reviewer synthesized **winning Branch B (0.6048)** with modeling ideas from lower-scoring Branch C (0.6039) for pairwise ranking and rank normalization, **yielding 0.6055536**.

Branch B+C Ideas → Synthesis 0.6055536

05 / EXPERIENCE —

Start your journey.

task.md

What to solve

Formal task contract: research objectives, dataset inputs, compute constraints, and official evaluation protocol.

- Dataset paths & split integrity
- Official sandbox evaluator

PERSONAL.md

How to work

User research preferences, model guidance, exploration style, tool permissions, and supervisory checkpoints.

- Model rotation (GPT / Gemini)
- Supervision & budget thresholds

SUPERVISION TIERS

Tier 1 Unattended Overnight Runs · **Tier 2** Checkpoint Review (META) · **Tier 3** Steer Working Priors


1. Define task.md → 2. Single-step test → 3. Autonomous search → 4. Parallel explore → 5. Inspect evidence

DECLARATIVE PROTOCOL

Human defines boundary and preferences (task.md + PERSONAL.md); SciOdyssey autonomously iterates, heals, and crystallizes verified evidence.

Autonomous by default.

Supervisable by design.



github.com/zc6600/research-agent-kuairand ↗

Team Good4AI: Chen Zhu · Zhou Ziyu · Shilin Xu · GE GAO · Jiran Li

Sources: slides/slides.md · docs/ARCHITECTURE.md · docs/FINAL_REPORT.md · KuaiRand-Pure submission ledger

TEAM GOOD4AI · ROBUST · SUPERVISABLE · BROAD SEARCH

A Persistent Research World for Autonomous Science

PROJECT WIKI & DEMO

sciodyyssey.zc6600.wiki ↗

Scan for live report & artifacts

