

AUTONOMOUS ML RESEARCH / TECHJAM 2026

SciOdyssey: Long-Horizon Autonomous ML Research Agent

Autonomous ML discovery through **pure evidence**.

PERSISTENT WORLD

FRESH SCIENTIST

EVIDENCE FIRST



One Agent. Any ML Task. End to End.

A general-purpose research agent that runs the entire research loop autonomously.

01 task.md

What to solve.

Objectives · constraints · evaluation

EXAMPLE

Rank short videos for each user and interaction context.

KuaiRand-Pure · long_view · GAUC + nDCG@5

02 PERSONAL.md

How you want it solved.

Preferences · research style · priorities


EXAMPLE

Run on macOS / Apple Silicon with uv.


≤ 15 min / experiment · never print credentials

Autonomous by default. Supervisable by design.


01 CLOSED-WORLD WORKFLOW



02 SINGLE TRAJECTORY



03 TREE SEARCH



DESIGN RESPONSE

Research is open-world, non-linear, and exploratory.

ADAPT

DIVERSIFY

LEARN

Preserve the
research world.

Reset the researcher.

Pass evidence, failures, and code to a fresh Scientist.

Two scientific roles. One safe runtime.

Scientist

Hypothesize · code · test · interpret

One trajectory

META

Audit · preserve · hand off

Across trajectories

Runtime

Isolation · cancellation · invariants

System lifetime

The Scientist chooses the science. META maintains the world. Runtime enforces mechanics, not scientific judgment.

Research is an **open-world** task. We need an **open-world** coding agent.

01 OPEN ACTION SPACE



Shell



Files



Browser



Git

Use what exists

02 SELF-RECOVERY



Read errors



Inspect docs



Inspect source



Install
packages

Fix what breaks

03 RUNTIME EXTENSION



MCP



Skills



CLI



Packages


Add what's missing

We chose a coding agent—not a predefined workflow around known tools.


Missing capability → inspect → extend environment → retry

The persistence boundary is explicit.

Environment




Research World




CONTEXT RESET

No previous reasoning trajectory

Fresh Scientist



META Review



Selective persistence / trajectory handoff

audit → scope → preserve

Runtime – deterministic mechanics



Workspace isolation



Runner + evaluator



Logging




Artifact capture



State operations

Each cycle leaves a world to build on.



PERSISTS metrics · code · failures · original reports

RESTARTS an independent reasoning trajectory

Keep facts distinct from intuition.

research world

- ├ evidence & provenance
- ├ verified knowledge
- ├ current research state
- ├ fallible priors
- └ State + implementation

Built to be re-examined.


Facts have evidence.

Priors can be overturned.

Code can be picked up again.

Original reports remain available for audit.

One starting point. Independent answers.



Explore in isolated worktrees.

Review after completion.

Autonomous research. Measured gains.

PRIMARY / RESEARCH AGENT

0.6059363

+0.0043363 absolute improvement

METRIC	Official FM	Ours
GAUC	0.6674000	0.6728421
nDCG@5	0.5357000	0.5390304
Primary	0.6016000	0.6059363

Primary = mean(GAUC, nDCG@5). Public validation only; hidden-test scoring is organizer-controlled.

Final recipe: 46 categorical features × 8-seed FM ensemble · NumPy / CPU

We are NOT tokenmaxxing.

Better research decisions, not more resources.

~\$10

from first line of code to final result

coding • debugging • agent runs • all experiments

MODEL

One GPT Plus + One Gemini Pro for 7 days

Two consumer subscriptions covered the entire research loop.


COMPUTE

1 × MacBook M2

All development and experiments ran on a single consumer laptop.

Performance came from the research loop – not a bigger budget.

The result stays connected to the evidence.



An audit trail from experiment to output.

4

Autonomous cycles

Within a 50-iteration cap

13

Named experiments

E001-E013 · 7 Full evaluations

0

Manual scientific interventions

After launch · 0 GPU-hours

170,588 prediction rows passed the unchanged Starter Kit alignment checker.

LLM tokens: 48.24M including cache reads; 4.02M excluding cache reads.

A retained state keeps the evidence inspectable.

The screenshot shows a 'Research Dashboard' with a 'RECORDED VALIDATION' status. The main content area displays a 'Retained Validation Result' for checkpoint 'S004', which is highlighted as a 'RETAINED STATE'. A green box highlights the 'SCORE-BEARING CHECKPOINT' section, which includes the checkpoint ID 'S004' and a description: '8-seed ensemble of 46-field Factorization Machine rankers incorporating leak-free historical user preferences (top tag, top tab, long-view rate bucket, interaction volume bucket, real-time tag/tab matches), secondary video tags, upload age, and per-seed Medium early stopping'. Below this, a table of metrics is shown: GAUC (0.672842), nDCG@5 (0.539030), Primary (0.605936), and Delta vs baseline (+0.004336). The evidence path is listed as 'system/evidence/cycle4-fm-rich46-ensemble8-full.json'. At the bottom, there are three buttons: 'Research Intuition', 'Recent Activity', and 'Run Summary'.

GAUC	nDCG@5	Primary	Delta vs baseline
0.672842	0.539030	0.605936	+0.004336

Evidence · system/evidence/cycle4-fm-rich46-ensemble8-full.json

Research Intuition Recent Activity Run Summary

Inspect the retained state and the evidence behind it.

Retained validation checkpoint · evidence path · reproducible implementation.

Reproducible evidence. Open questions.

- | | | |
|----|--|--------------------|
| 01 | One task does not establish generality. | External validity |
| 02 | META, State, and reset are not isolated. | Causal attribution |
| 03 | Summaries can lose context and anchor. | Memory quality |
-

Next questions: broader tasks, targeted ablations, and less repeated execution.

Five people. One autonomous research loop.

OVERLAPPING RESPONSIBILITIES / SHARED WORLD



Chen Zhu

TEAM LEAD / SYSTEM

Direction · architecture ·
integration



Zhou Ziyu

MODELING / EXPERIMENTS

Features · ensembles ·
evidence review



Shilin Xu

DATA / METRICS

Data workflow · contracts ·
alignment



GE GAO

RUNTIME / RELIABILITY

Execution · testing ·
integration support



Jiran Li

EVIDENCE /
REPRODUCIBILITY

Telemetry ·
documentation ·
packaging

Keep the experience. Reopen the search.

```
# From the configured repo, run one autonomous cycle
$ ./scripts/research-agent step --cli codex \
  --target /absolute/path/to/project --allow-edits
```

github.com/zc6600/research-agent-kuairand ↗

Code · Technical report · Experiment ledger · Checked output

The run taught us more than the final score.

What survived the experiments – and what remains unsettled.

01 Diversity is a prior.

Keep `gemini-3.7-flash` as META while Scientist rotates across `gpt-5.6-sol`, `gemini-3.7-flash`, and `gpt-5.6-luna`. This widens priors; it does not prove the gain.

02 A score is not a claim.

Audits found leakage and sampling mismatches in competing experiments. META can scope evidence; it is not a formal verifier.

03 Failure can be evidence.

After evaluation, NumPy `float32` values broke serialization. The Scientist fixed the writer, reran, and kept the measurements.

04 Reset $\neq$ no anchor.

Fresh cognition reopens the search, but shared summaries can still omit context or anchor future work. Share evidence later than cognition.

OBSERVED, NOT PROVEN

Parallel/Synthesis reached public-validation Primary **0.6055536**; the canonical E001-E013 frontier remains **0.6059363**.

4 cycles


0 post-launch interventions

0 GPU-hours

The difference is not orchestration. It is recoverability.

LANGGRAPH-STYLE ORCHESTRATION

predefined graph




- ⊘ API fails
- ↓
- ⊘ no recovery tool
- ↓
- ⊘ **DEAD END**

Capabilities are bounded by what developers anticipated.

CODING-AGENT HARNESS

open action space



fail → inspect → learn → modify environment → retry

The agent can acquire the missing capability at runtime.

A clean repo. A continuous agent.

FIELD NOTE / 01

Two practical patterns for keeping autonomous work moving – across branches and across time.

01 MANY AGENTS



Give every agent a branch.

main

clean trunk

● agent A / branch

● agent B / branch

● agent C / branch

review

→ merge

Let agents work independently. Conflicts stay local, and the main branch stays clean.

02 NO QUOTA / NO LAPTOP



Leave the next move in GitHub.

01

**Issue /
prompt**

inject the task



02

**GitHub
Action**

run the agent



03

Commit

review the change

Connect the agent to GitHub Actions. It can code and commit while you are offline; review before merging.