

AUTONOMOUS ML RESEARCH / TECHJAM 2026

# SciOdyssey

## Long-Horizon Autonomous ML Research Agent

A persistent research world. A fresh Scientist every cycle.

PERSISTENT WORLD

FRESH SCIENTIST

EVIDENCE FIRST

Problem → System → Experience → Evaluation → Insights



# Can an agent carry a research task to completion?

Turn a task contract into a tested model, a reproducible implementation, and evidence another researcher can inspect.

---

## Explore

Choose hypotheses, write code, and run experiments.

Scientific judgment

---

## Recover

Handle failures and revise unsuccessful ideas.

Long-running work

---

## Deliver


Retain a valid model and the evidence behind it.

Verifiable progress

Our test case: short-video ranking on KuaiRand-Pure, with a fixed evaluator and public-validation boundary.

# Why research agents fall short

## 01 WORKFLOW




 **ERROR**  
↓  
dead end

### Fixed recovery paths

Unexpected errors need human help.

## 02 SINGLE AGENT



### History shapes search

Errors can persist

## 03 TREE SEARCH




### Repeated context reads


Each small edit pays the reading cost.

# A system for sustained research


## Environment



## Research World



## Fresh Scientist



## META Review



Selective persistence / trajectory handoff

audit → scope → preserve

## Coding-agent harness + runtime



Workspace isolation



Runner + evaluator



Logging




Artifact capture



State operations

# Each cycle leaves a world to build on.



**PERSISTS** metrics · code · failures · original reports

**RESTARTS** an independent reasoning trajectory

# Start with the task and working constraints.

The operator defines the objective and constraints; the agent chooses and executes the experiments.

---

01 task.md

---

## What to solve.

Objectives · constraints · evaluation

EXAMPLE

Rank short videos for each user and interaction context.

KuaiRand-Pure · long\_view · GAUC + nDCG@5

---

02 PERSONAL.md

---

## How you want it solved.

Preferences · research style · priorities

EXAMPLE

Run on macOS / Apple Silicon with uv.

≤ 15 min / experiment · never print credentials

---

Autonomous by default. Supervisable by design.

# Evidence changed the next experiment.

E001-E013: reject weak ideas, improve representations, then test ensembles.

|      |                                                           |                  |
|------|-----------------------------------------------------------|------------------|
| E001 | Screen historical target encodings; reject weak variants. | Medium screening |
| E003 | Establish a valid rich-FM checkpoint.                     | Full · 0.6016310 |
| E008 | Retain a five-seed, 38-field FM ensemble.                 | Full · 0.6040901 |
| E013 | Select the eight-seed, 46-field FM ensemble.              | Full · 0.6059363 |

These are milestones within four autonomous cycles. Medium screens ideas; Full evaluations support retained score claims.



# An evaluation failure became a recoverable step.


Cycle 1 completed training and evaluation, then failed while writing the evidence file.

|         |                                                   |                    |
|---------|---------------------------------------------------|--------------------|
| Failure | NumPy float32 values broke JSON serialization.    | Evidence writer    |
| Repair  | The Scientist converted scalars and reran.        | Agent-authored fix |
| Retain  | Printed measurements and saved evidence survived. | Auditable recovery |

No human selected this repair or edited the code. This is one observed recovery case, not a guarantee for every failure.



# A retained state keeps the evidence inspectable.



Inspect the retained state and the evidence behind it.

Retained validation checkpoint · evidence path · reproducible implementation.

# Autonomous research. Measured gains.

PRIMARY / RESEARCH AGENT

0.6059363

+0.0043363 absolute improvement

| METRIC  | Official FM | Ours      |
|---------|-------------|-----------|
| GAUC    | 0.6674000   | 0.6728421 |
| nDCG@5  | 0.5357000   | 0.5390304 |
| Primary | 0.6016000   | 0.6059363 |

Primary = mean(GAUC, nDCG@5). Public validation only; hidden-test scoring is organizer-controlled.

Final recipe: 46 categorical features × 8-seed FM ensemble · NumPy / CPU

# An audit trail from experiment to output.

4

**Autonomous cycles**

Within a 50-iteration cap

13

**Named experiments**

E001-E013 · 7 Full evaluations

0

**Manual scientific interventions**

After launch · 0 GPU-hours

170,588 prediction rows passed the unchanged Starter Kit alignment checker.

LLM tokens: 48.24M including cache reads; 4.02M excluding cache reads.

# Resource use is part of the evidence.

Project-level subscription estimate, with run-level telemetry reported separately.

~\$10

from first line of code to final result

coding · debugging · agent runs · all experiments

---

## MODEL

One GPT Plus + One Gemini Pro for 7 days

Two consumer subscriptions covered the entire research loop.

---


## COMPUTE

1 × MacBook M2

All development and experiments ran on a single consumer laptop.

The recorded runs used CPU training; LLM usage remains a separate cost.

# Compare scores alongside their evidence.



Runs differ in model, budget, and evidence quality; this is not a controlled causal comparison.

# The run taught us more than the final score.

What survived the experiments – and what remains unsettled.

## 01 Diversity is a prior.

Keep `gemini-3.7-flash` as META while Scientist rotates across `gpt-5.6-sol`, `gemini-3.7-flash`, and `gpt-5.6-luna`. This widens priors; it does not prove the gain.

## 02 A score is not a claim.

Audits found leakage and sampling mismatches in competing experiments. META can scope evidence; it is not a formal verifier.

## 03 Failure can be evidence.

After evaluation, NumPy `float32` values broke serialization. The Scientist fixed the writer, reran, and kept the measurements.

## 04 Reset $\neq$ no anchor.

Fresh cognition reopens the search, but shared summaries can still omit context or anchor future work. Share evidence later than cognition.

### OBSERVED, NOT PROVEN

Parallel/Synthesis reached public-validation Primary 0.6055536; the canonical E001-E013 frontier remains 0.6059363.

4 cycles

0 post-launch interventions

0 GPU-hours

# Reproducible evidence. Open questions.

- 
- |    |                                         |                   |
|----|-----------------------------------------|-------------------|
| 01 | One task does not establish generality. | External validity |
|----|-----------------------------------------|-------------------|
- 
- |    |                                          |                    |
|----|------------------------------------------|--------------------|
| 02 | META, State, and reset are not isolated. | Causal attribution |
|----|------------------------------------------|--------------------|
- 
- |    |                                        |                |
|----|----------------------------------------|----------------|
| 03 | Summaries can lose context and anchor. | Memory quality |
|----|----------------------------------------|----------------|
- 
- Next questions: broader tasks, targeted ablations, and less repeated execution.



# A clean repo. A continuous agent.

FIELD NOTE / 01

Engineering practices to apply next: isolate parallel changes and make unattended work reviewable.

## 01 MANY AGENTS



### Give every agent a branch.

**main**

clean trunk

● agent A / branch

● agent B / branch

● agent C / branch

**review**

→ merge

Let agents work independently. Conflicts stay local, and the main branch stays clean.

## 02 NO QUOTA / NO LAPTOP



### Leave the next move in GitHub.

01

**Issue /  
prompt**

inject the task



02

**GitHub  
Action**

run the agent



03

**Commit**review the  
change

Connect the agent to GitHub Actions. It can code and commit while you are offline; review before merging.

# Keep the experience. Reopen the search.

---

```
# From the configured repo, run one autonomous cycle
$ ./scripts/research-agent step --cli codex \
  --target /absolute/path/to/project --allow-edits
```

---

[github.com/zc6600/research-agent-kuairand](https://github.com/zc6600/research-agent-kuairand) ↗

Code · Technical report · Experiment ledger · Checked output

Research is an **open-world** task.  
We need an **open-world** coding agent.

### 01 OPEN ACTION SPACE



Shell



Files



Browser



Git

Use what exists

### 02 SELF-RECOVERY



Read errors



Inspect docs



Inspect  
source



Install  
packages

Fix what breaks

### 03 RUNTIME EXTENSION



MCP



Skills



CLI



Packages

Add what's missing

WORKFLOW EXAMPLE • LANGGRAPH

**LangGraph** makes a known workflow explicit as nodes + edges.

Missing capability? Leave the graph → inspect → extend → retry

# Five people. One autonomous research loop.

OVERLAPPING RESPONSIBILITIES / SHARED WORLD



Chen Zhu

TEAM LEAD / SYSTEM

Direction · architecture ·  
integration



Zhou Ziyu

MODELING / EXPERIMENTS

Features · ensembles ·  
evidence review



Shilin Xu

DATA / METRICS

Data workflow · contracts  
· alignment



GE GAO

RUNTIME / RELIABILITY

Execution · testing ·  
integration support



Jiran Li


EVIDENCE /  
REPRODUCIBILITY

Telemetry · documentation  
· packaging

Research is an **open-world** task.  
We need an **open-world** coding agent.

## LANGGRAPH-STYLE ORCHESTRATION


predefined graph



Capabilities are bounded by what developers anticipated.

## CODING-AGENT HARNESS

open action space



fail → inspect → learn → modify environment → retry

The agent can acquire the missing capability at runtime.