

01 / PROBLEM

Different lifetimes. Shared evidence.

Long-lived reasoning can anchor the search; rebuilding context is costly. SciOdyssey retains the world while reopening judgment with a fresh Scientist.

Research environment

Task contract + research preferences
Data + evaluator + baseline · budget + permissions

02 / MOTIVATION



Research world

PERSISTENT / ON DISK

EVIDENCE

Knowledge + experiment ledger
Metrics, failures and original reports

CURATED

Research brief + fallible priors
Current understanding; open to revision

CODE

Retained implementation + State
Recoverable version with evidence and provenance

↓ CONTEXT RESET Inherit experience, not prior reasoning.

03 / SYSTEM




Scientist

SCIENTIFIC JUDGMENT

- Read the research world
- Form hypotheses
- Code, train and evaluate
- Interpret + report evidence

Chooses the science within one trajectory.



META

PERSISTENCE AUTHORITY

- Audit claims ↔ evidence
- Scope conclusions
- Curate shared memory
- Preserve optional State

Maintains what survives across trajectories.

Audit → scope → preserve → fresh Scientist

Runtime / deterministic mechanics

Isolation · cancellation · runner + evaluator
Logging · artifact capture · State operations

One world. Several independent searches.



Isolated worktrees, then post-hoc review. Optional synthesis keeps one implementation parent and reference evidence; code and memory are not auto-merged.

Framework capability, not an isolated cause of benchmark gain.

04 / EXPERIENCE

One agent. Any ML task. End to end.

task.md

PERSONAL.md

What to solve

Objectives · constraints · evaluation

How you want it solved

Preferences · research style · priorities


Define the task → Delegate research → Review evidence

05 / EVALUATION

Research outcomes, in context.

KuaiRand-Pure · public validation · one recorded run per condition

Primary score ↑



* Codex is provisional. Gemini = gemini-3.7-flash. The Cycle-2 interval is a token-accounting bound. Models and protocols differ: no causal, significance or token-efficiency claim is made.

Verified submission · public validation

Metric	Official FM	SciOdyssey	Absolute Δ
Primary	0.6016000	0.6059363	+0.0043363
GAUC	0.6674000	0.6728421	+0.0054421
nDCG@5	0.5357000	0.5390304	+0.0033304

Primary = mean(GAUC, nDCG@5). Final recipe: 46 categorical features × eight-seed FM ensemble · NumPy / CPU.

4 autonomous cycles 13 named experiments 0 manual scientific interventions*

7 Full evaluations · 0 GPU-hours · 170,588 prediction rows passed the unchanged Starter Kit alignment checker.

* After launch, in the retained submission record.

06 / INSIGHTS

Observed behavior, not just a score.

Recovery

After evaluation, NumPy float32 values broke serialization. The Scientist repaired the writer and retained the evidence.

Audit

Reviews found leakage and sampling mismatches. META scopes evidence; it is not a formal verifier.

Open questions. One task does not establish generality. META, State and reset effects are not isolated; summaries can still lose context and anchor future work.

Autonomous by default. Supervisable by design.

Coding agent

EXAMPLE INPUT

Use the research-agent skill for task.md. Follow PERSONAL.md. Run autonomously. Let me review the results and evidence.